



1. a)

**Plot of gibbs sampled distribution for my**



**Plot of gibbs sampled distribution for sigma**



Inefficiency factor my: 0.52

Inefficiency factor sigma: 1.05

**Conclusion:** As the plots suggest, we can't distinguish any convergence for the mean and standard deviation as we clearly can see that the values oscillate around the means (red line). Moreover, the inefficiency factors shows how correlated the draws are

## Code:

```
data=readRDS("Precipitation.rds") #Reading data
data = log(data)

#a)

# Initial setting (prior settings)
my0 = 1
tao0=1
v0=1
sigma0sq=1
nDraws = 1000
n = nDraws
true_mean = mean(data)
vn = v0 + n

#Lists for storing draws
my_vals = rep(0,nDraws)
sigma_vals = rep(0,nDraws)
sigma_vals[1] = 1 #Initial sigmasq value to be used

#Drawing posteriors
for(i in 1:nDraws) {

  w = (n/sigma_vals[i])/((n/sigma_vals[i])+(1/tao0**2)) #Calculating weights
  myn = w*true_mean+(1-w)*my0 #Calculating myn
  taonsq = 1/((n/sigma_vals[i])+(1/tao0)) #Calculating taonsq
  my_vals[i] = rnorm(1,mean=myn, taonsq) #Calculating posterior my_val

  if (i <= (nDraws-1)) {
    chi_val = rchisq(1,vn) #Drawing random value from chi distribution
    sigmansq = (v0*sigma0sq+sum((data-my_vals[i])**2))/(n+v0) #Calculating approximation
    sigmasq
    sigma_vals[i+1] = vn*sigmansq/chi_val #Calculating posterior sigma_val
  }
}
```

#Plotting my convergence

```
plot(my_vals[10:nDraws], .main = "Plot of gibbs sampled distribution for my",type = "l")
abline(h=mean(my_vals), col="red") #Expected value
```

#Plotting sigma convergence

```
plot(sqrt(sigma_vals[10:nDraws]), .main = "Plot of gibbs sampled distribution for sigma",type
```

```
= "I")  
abline(h=mean(sqrt(sigma_vals)), col="red") #Expected value
```

#Calculating inefficiency factor (IF)

#Autocorrelation for my

```
a_my = acf(my_vals)
```

```
inefficiency_factor_my = 1+2*sum(a_my$acf[-1])
```

```
inefficiency_factor_my
```

```
## [1] 0.5193648
```

#Autocorrelation for my

```
a_sigma = acf(sqrt(sigma_vals))
```

```
inefficiency_factor_sigma = 1+2*sum(a_sigma$acf[-1])
```

```
inefficiency_factor_sigma
```

```
## [1] 1.04877
```



#Conclusion: As the plots suggest, we can't distinguish any convergence for the mean and standard deviation

#as we clearly can see that the values oscillate around the means (red line)

#Moreover, the inefficiency factors shows how correlated the draws are, as we can see, the results are

#relatively close to 1 which indicates low correlation

b)



**Conclusion:** As we can see, the draws from the predictions differ slightly from the true distribution. The main difference is that the normal distribution have a bit longer tails on each end, however the distributions seem to be centred around the same values in both plots

**Code:**

#b)

#1) Histogram or kernel density estimate of the daily preception  $y_1, \dots, y_n$

```
hist(data) #Histogram with true data
```

```
#2) The resulting posterior predictive density  $p(\tilde{y}|y_1, \dots, y_n)$  using the simulated posterior draws from (a)
```

```
predictions = rep(0,nDraws)
```

```
for (i in 1:nDraws) {
```

```
predictions[i] = rnorm(1,my_vals[i],sigma_vals[i])
```

```
}
```

```
hist(predictions)
```

```
#Conclusion: As we can see, the draws from the predictions differ slightly from the true distribution.
```

```
#The main difference is that the normal distribution have a bit longer tails on each end, however
```

```
#the distributions seem to be centered around the same values in both plots
```

2. a)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	1.07244	0.03077	34.848	< 2e-16	***
PowerSeller	-0.02054	0.03678	-0.558	0.5765	
VerifyID	-0.39452	0.09243	-4.268	1.97e-05	***
Sealed	0.44384	0.05056	8.778	< 2e-16	***
Minblem	-0.05220	0.06020	-0.867	0.3859	
MajBlem	-0.22087	0.09144	-2.416	0.0157	*
LargNeg	0.07067	0.05633	1.255	0.2096	
LogBook	-0.12068	0.02896	-4.166	3.09e-05	***
MinBidShare	-1.89410	0.07124	-26.588	< 2e-16	***

the z-value shows that VerifyID, Sealed, MajBlem, LogBook and MinBidShare are significant

**Code:**

```
# (a) Obtain the maximum likelihood estimator of  $\beta$  in the Poisson regression model for the eBay data
```

```
# [Hint: glm.R, don't forget that glm() adds its own intercept so don't input the covariate Const].
```

```
# Which covariates are significant?
```

```
library(mvtnorm)
```

```
ebay_data <- (read.table("eBayNumberOfBidderData.dat", header=TRUE))
```

```
Xnames <- names(ebay_data[,2:ncol(ebay_data)]) #Names of columns
```

```
n_obs <- dim(ebay_data)[1] #number of observations
```

```
data <- ebay_data[, -2] #training and test data without Const
```

```
#y <- as.numeric(data[, 1]) #target output
```

```
# maximum likelihood estimator of beta
```

```
model = glm(formula=nBids~., data=data, family=poisson())
```

```
summary(model)
```

```
##
```

```
## Call:
```

```
## glm(formula = nBids ~ ., family = poisson(), data = data)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -3.5800 -0.7222 -0.0441  0.5269  2.4605
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept) 1.07244      0.03077 34.848 < 2e-16 ***
## PowerSeller -0.02054     0.03678 -0.558  0.5765
## VerifyID    -0.39452     0.09243 -4.268 1.97e-05 ***
## Sealed      0.44384      0.05056  8.778 < 2e-16 ***
## Minblem     -0.05220     0.06020 -0.867  0.3859
## MajBlem     -0.22087     0.09144 -2.416  0.0157 *
## LargNeg     0.07067      0.05633  1.255  0.2096
## LogBook     -0.12068     0.02896 -4.166 3.09e-05 ***
## MinBidShare -1.89410     0.07124 -26.588 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 2151.28  on 999  degrees of freedom
## Residual deviance:  867.47  on 991  degrees of freedom
## AIC: 3610.3
##
## Number of Fisher Scoring iterations: 5
```

*the z-value shows that VerifyID, Sealed, MajBlem, LogBook and MinBidShare are significant*

b)

The posterior mode:

```
[1,] 1.06984118 "Const"
[2,] -0.02051246 "PowerSeller"
[3,] -0.39300599 "VerifyID"
[4,] 0.44355549 "Sealed"
[5,] -0.05246627 "Minblem"
[6,] -0.22123840 "MajBlem"
[7,] 0.07069683 "LargNeg"
[8,] -0.12021767 "LogBook"
[9,] -1.89198501 "MinBidShare"
```

Posterior standard deviation:

```
Const      PowerSeller  VerifyID    Sealed      Minblem    MajBlem    LargNeg
0.03074837 0.03678418 0.09227871 0.05057448 0.06020470 0.09146070 0.05634767
LogBook    MinBidShare
0.02895635 0.07109682
```

the result seems reasonable since it is similar to a)

**Code:**

```
# (b) Let's do a Bayesian analysis of the Poisson regression. Let the prior be  $\beta \sim N[0, 100 \cdot (X^T X)^{-1}]$ 
# where  $X$  is the  $n \times p$  covariate matrix. This is a commonly used prior, which is called Zellner's g-prior.
# Assume first that the posterior density is approximately multivariate normal:
#  $\beta|y \sim N(\tilde{\beta}, Jy^{-1}(\tilde{\beta}))$ 
# where  $\tilde{\beta}$  is the posterior mode and  $Jy(\tilde{\beta})$  is the negative Hessian at the posterior mode.
#  $\tilde{\beta}$  and  $Jy(\tilde{\beta})$  can be obtained by numerical optimization (optim.R)
# exactly like you already did for the logistic regression in Lab 2
# (but with the log posterior function replaced by the corresponding one for the Poisson model,
# which you have to code up.).
```

```
#function for Poisson logPosterior
```

```
logPostPoisson = function(beta,y,X,mu,Sigma) {
  beta <- matrix(beta, nrow=1)
  linPred <- X%*%t(beta)
  logLik <- sum( - exp(linPred) + linPred*y - log(factorial(y)) ) #loglikelihood
  logPrior <- dmvnorm(beta, mu, Sigma, log=TRUE); #prior
  return(logLik + logPrior) #prior and likelihood are added since we use log
}
```

```
#initialize variables
```

```
y <- as.numeric(ebay_data[,1]) #target output
X <- as.matrix(ebay_data[,-1]) #training data
X <- apply(X, c(1, 2), as.numeric) #convert from characters to doubles
n_col <- dim(X)[2] #number of columns in training data
initVal <- rep(0,n_col) #initial values for Beta
mu <- as.matrix(rep(0,n_col)) # Prior mean vector for Beta
Sigma <- 100*solve(t(X)%*%X) # Prior covariance matrix for Beta
logPost <- logPostPoisson; #function to be minimized
```

```
set.seed(123)
```

```
OptimRes <-
optim(initVal,logPost,gr=NULL,y,X,mu,Sigma,method=c("BFGS"),control=list(fnscale=-1),hessian=TRUE)
```

```
# Printing the results of mean and variance of normally distributed Beta
```

```
OptimBeta <- OptimRes$par
names(OptimBeta) <- Xnames # Naming the coefficient by covariates
PostCov <- solve(-OptimRes$hessian) #get covariance matrix by inverting the negative hessian
approxPostStd <- sqrt(diag(PostCov)) # Computing approximate standard deviations.
names(approxPostStd) <- Xnames # Naming the coefficient by covariates
print('Posterior beta values:')
```



```
## [1] "Posterior beta values:"
```

```
print(OptimBeta)
```

```
##      Const PowerSeller  VerifyID      Sealed      Minblem      MajBlem  
## 1.06984118 -0.02051246 -0.39300599 0.44355549 -0.05246627 -0.22123840  
##      LargNeg      LogBook MinBidShare  
## 0.07069683 -0.12021767 -1.89198501
```

```
print('Posterior standard deviation:')
```

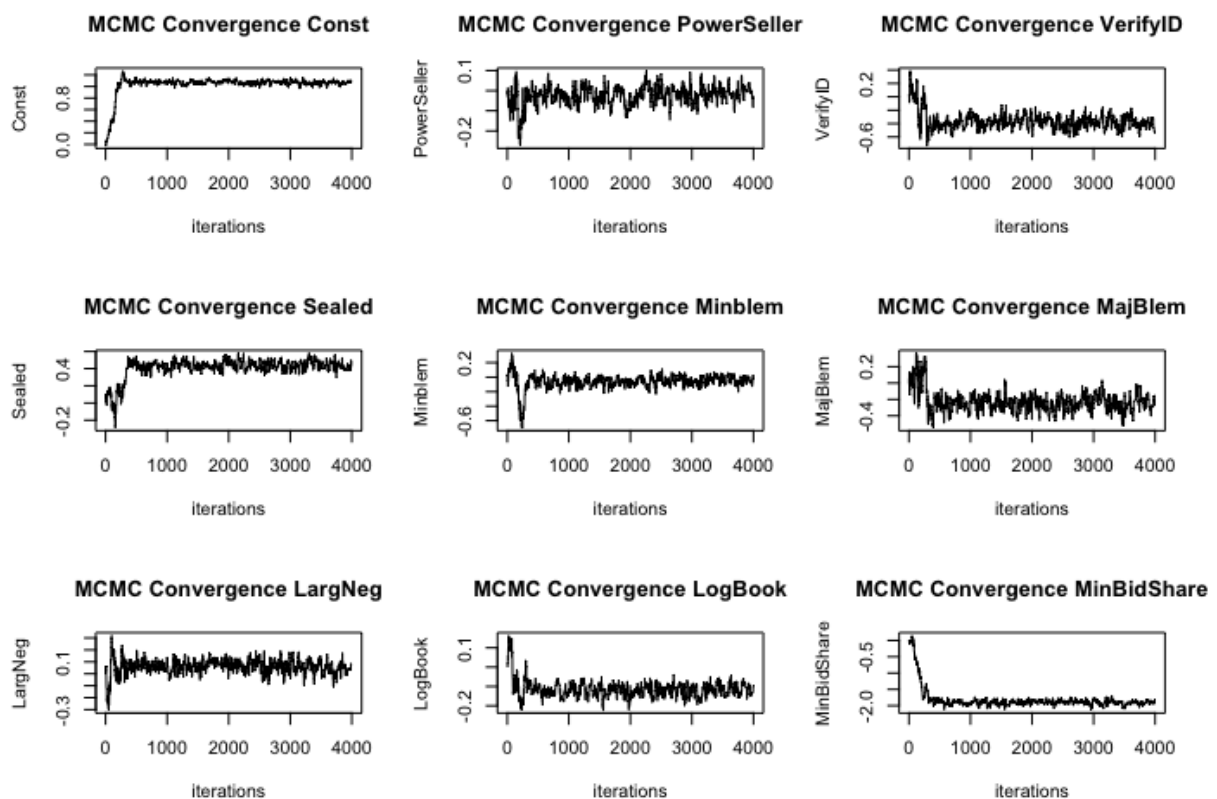
```
## [1] "Posterior standard deviation:"
```

```
print(approxPostStd)
```

```
##      Const PowerSeller  VerifyID      Sealed      Minblem      MajBlem  
## 0.03074837 0.03678418 0.09227871 0.05057448 0.06020470 0.09146070  
##      LargNeg      LogBook MinBidShare  
## 0.05634767 0.02895635 0.07109682
```

*#the result seems reasonable since they are similar to a)*

c)



Betavalues are converging as we can see in the plot.

average\_alpha = 0.2531, so we have average acceptance probability between 25-30% which means  $c=0.7$  was appropriate.

#### Code:

```
# (c) Let's simulate from the actual posterior of beta using the Metropolis algorithm  
# and compare the results with the approximate results in b).  
# Program a general function that uses the Metropolis algorithm to generate random draws  
from an  
# arbitrary posterior density. In order to show that it is a general function for  
# any model, we denote the vector of model parameters by  $\theta$ . Let the proposal density  
# be the multivariate normal density mentioned in Lecture 8 (random walk Metropolis):  
#  $\theta_p | \theta(i-1) \sim N(\theta(i-1), c \cdot \Sigma)$  where  $\Sigma = J_y^{-1}(\beta^*)$  was obtained in b).  
# The value c is a tuning parameter and should be an input to your Metropolis function.  
# The user of your Metropolis function should be able to supply her own posterior density  
function, not  
# necessarily for the Poisson regression, and still be able to use your Metropolis  
# function. This is not so straightforward, unless you have come across function  
# objects in R. The note HowToCodeRWM.pdf in Lisam describes how you can do this in R.
```

```
#Now, use your new Metropolis function to sample from the posterior of  $\beta$   
# in the Poisson regression for the eBay dataset. Assess MCMC convergence by graphical  
methods.
```

```
RWMSampler <- function(oldTheta, logPostFunc, c, PostCov, ... ) {  
  proposedTheta <- rmvnorm(1, mean=oldTheta, sigma=c*PostCov)  
  alpha <- min(1, exp(logPostFunc(proposedTheta, ...) - logPostFunc(oldTheta, ...)) )  
  t <- runif(1)  
  if (alpha>t) {  
    return(list(proposedTheta, alpha))  
  } else {  
    return(list(oldTheta, alpha))  
  }  
}
```

```
nDraws <- 4000
```

```
c <- 0.7
```

```
beta <- matrix(0, nDraws, n_col) #matrix to fill with samples from RWM
```

```
alpha_vector <- rep(0, nDraws) #will be used to calculate mean alpha
```

```
logPost <- logPostPoisson
```

```
for(i in 1:nDraws) {  
  output <- RWMSampler(beta[i,], logPost, c, PostCov, y, X, mu, Sigma)  
  alpha_vector[i] <- output[[2]]  
  if(i<nDraws) {  
    beta[i+1,] <- output[[1]]  
  }  
}
```

```

}

iterations=seq(1,nDraws,1)
par(mfrow=c(3,3))
for (i in 1:n_col) {
  plot(iterations, beta[,i], type="l", main=paste("MCMC Convergence", Xnames[i]),
        ylab=Xnames[i])
}

par(mfrow=c(1,1), new=FALSE)

# Calculate average alpha
average_alpha <- mean(alpha_vector)
average_alpha #0.2531, has average acceptance probability between 25-30%

## [1] 0.2621183

```

d)



Calculations give 74% chance that we have no bidders with the given characteristics.

**Code:**

```

# (d) Use the MCMC draws from c) to simulate from the predictive distribution of
# the number of bidders in a new auction with the characteristics below.
#Plot the predictive distribution. What is the probability of no bidders in this new auction?
# PowerSeller = 1

```

```

# VerifyID = 0
# Sealed = 1
# MinBlem = 0
# MajBlem = 1
# LargNeg = 0
# LogBook = 1.2
# MinBidShare = 0.8

nDraws <- 10000
x <- c(1,1,0,1,0,1,0,1.2,0.8)
post_beta <- beta[(1000:nrow(beta)),] #select samples after convergence
mean <- exp(post_beta%*%x) #vector of means
samples <- rpois(nDraws, mean)
barplot(table(samples), main="Histogram of number of bidders")

prob_no_bidders <- sum(samples==0)/sum(samples) #probability of no bidders
print(prob_no_bidders)

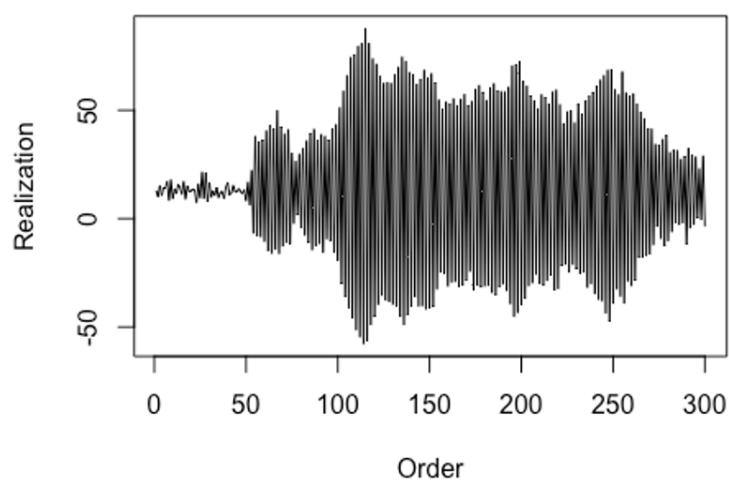
## [1] 0.739524

#the probability of no bidders is 74%


```

3. a)


**Realization for  $\phi = -1$**




**Realization for  $\phi = -0.5$**



**Realization for  $\phi = 0.5$**



**Realization for  $\phi = 1$**



**Conclusion:** Increased  $\phi$  values reduces oscillations. We can see clear oscillations for  $\phi = -1$  and that the output becomes more and more correlated with increased  $\phi$  values

**Code:**

```
#a)
```

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```

## Loading required package: ggplot2

## rstan (Version 2.21.8, GitRev: 2e1f913d3ca3)

## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)

# Initial values
my=13
sigmasq=3
T=300
x1 = my
phi = seq(-1,1,0.25)

#Autoregression function
ar_function = function(my,sigmasq,T,x1,phi){
  result = rep(0,T)
  result[1] = x1
  for(i in 2:T) {
    epsilon = rnorm(1,0,sigmasq)
    result[i] = my + phi*(result[i-1]-my) + epsilon
  }
  return(result)
}

#Calculating output
result = matrix(0,T,length(phi))
for(i in 1:length(phi)) {
  result[,i] = ar_function(my,sigmasq,T,x1,phi[i])
}

#Plotting result for phi=-1 and phi=1
plot(seq(1,300,1),result[,1],xlab="Order",ylab="Realization",main="Realization for phi =
-1",type="l")

plot(seq(1,300,1),result[,3],xlab="Order",ylab="Realization",main="Realization for phi =
-0.5",type="l")

plot(seq(1,300,1),result[,7],xlab="Order",ylab="Realization",main="Realization for phi =
0.5",type="l")

plot(seq(1,300,1),result[,9],xlab="Order",ylab="Realization",main="Realization for phi =
1",type="l")

#Conclusion: Increased phi values reduces oscilations. We can see clear oscilations for
phi=-1
#and that the output becomes more and more correlated with increased phi values

```

b)


*Table 1: Values for x model*

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	12.84	0.00	0.20	12.43	12.71	12.84	12.97	13.24	3440	1
phi	0.12	0.00	0.06	0.01	0.08	0.12	0.16	0.24	3620	1
sigma2	9.02	0.01	0.76	7.66	8.49	8.99	9.51	10.61	3611	1
lp__	-474.33	0.03	1.30	-477.78	-474.90	-473.98	-473.39	-472.89	1970	1

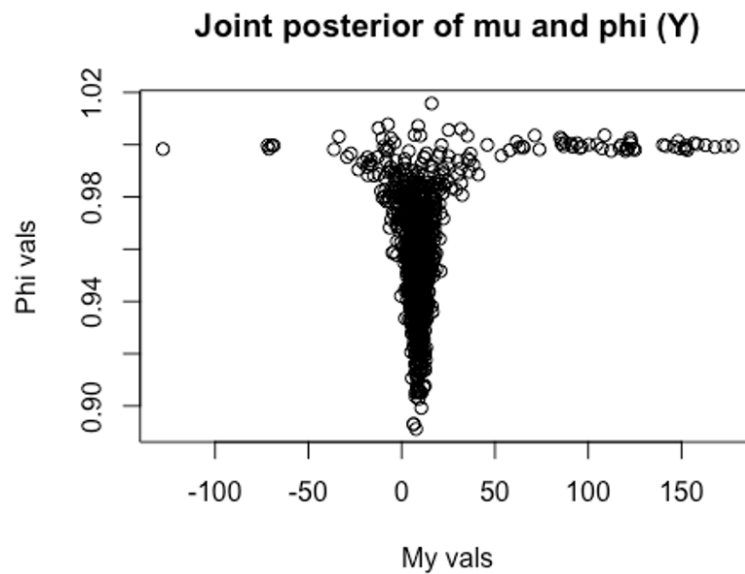
*Table 2: Values for y model*

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	14.76	1.33	24.52	-40.56	9.38	14.82	19.58	76.75	342	1.01
phi	0.98	0.00	0.02	0.94	0.96	0.98	0.99	1.00	779	1.01
sigma2	8.64	0.02	0.72	7.37	8.14	8.58	9.11	10.20	1273	1.00
lp__	-467.64	0.05	1.38	-470.83	-468.53	-467.35	-466.51	-465.89	764	1.01

**Conclusion:** As the values in the credible interval are more narrow for the one where  $\phi=0.2$  were used, we get a better understanding of the true values for this one compared to the one where  $\phi=0.95$  were used. For this one, the values are really wide in the credible intervals







**Conclusion:** We can distinguish a convergence for the first sampler. However it is more difficult to distinguish a convergence for the second sampler as shown in the figure. This follows the previous reasoning, that the second sampler have a wider span for its values compared to the first sampler

**Code:**

```
#b)
#Drawing x1:T values using phi=0.2
x=rep(0,T)
x=ar_function(my,sigmasq,T,x1,0.2)

#Drawing y1:T values using phi=0.95
y=rep(0,T)
y=ar_function(my,sigmasq,T,x1,0.95)

#Implementing Stan model

StanModel = '
data {
  int<lower=0> N;
  vector[N] z;
}
parameters {
  real mu;
  real phi;
  real<lower=0> sigma2;
}
model {
```

```

for(i in 2:N){
  z[i] ~ normal(mu+phi*(z[i-1]-mu),sqrt(sigma2));
}
}

#For X
data <- list(N=T, z=x)
fit_x <- stan(model_code=StanModel,data=data)

#For Y
data <- list(N=T, z=y)
fit_y <- stan(model_code=StanModel,data=data)

#Extract posterior draws
post_x = extract(fit_x)
post_y = extract(fit_y)

#i)
#AR-process for X
#Calculating mean and credible interval for all parameters
#Can be achived with print(fit_x) too
mean_mu = mean(post_x$mu)
cred_int_mu = quantile(post_x$mu,probs=c(0.025,0.975))
print(mean_mu)

## [1] 12.79963

print(cred_int_mu)

## 2.5% 97.5%
## 12.43112 13.17352

mean_phi= mean(post_x$phi)
cred_int_phi = quantile(post_x$phi,probs=c(0.025,0.975))
print(mean_phi)

## [1] 0.1432679

print(cred_int_phi)

## 2.5% 97.5%
## 0.03507495 0.25712395

mean_sigma2 = mean(post_x$sigma2)
cred_int_sigma2 = quantile(post_x$sigma2,probs=c(0.025,0.975))
print(mean_sigma2)

## [1] 7.918592

print(cred_int_sigma2)

```

```

## 2.5% 97.5%
## 6.726024 9.330297

#AR-process for Y
#Calculating mean and credible interval for all parameters
mean_mu = mean(post_y$mu)
cred_int_mu = quantile(post_y$mu,probs=c(0.025,0.975))
print(mean_mu)

## [1] 10.04247

print(cred_int_mu)

## 2.5% 97.5%
## -16.08935 50.19349

mean_phi = mean(post_y$phi)
cred_int_phi = quantile(post_y$phi,probs=c(0.025,0.975))
print(mean_phi)

## [1] 0.957765

print(cred_int_phi)

## 2.5% 97.5%
## 0.9116279 1.0007403

mean_sigma2 = mean(post_y$sigma2)
cred_int_sigma2 = quantile(post_y$sigma2,probs=c(0.025,0.975))
print(mean_sigma2)

## [1] 9.138192

print(cred_int_sigma2)

## 2.5% 97.5%
## 7.744661 10.745055

#Plotting using print
print(fit_x)

## Inference for Stan model: d85eaf939f1fe064d5bdde49c5af5b23.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##      mean se_mean sd 2.5% 25% 50% 75% 97.5% n_eff Rhat
## mu    12.80  0.00 0.19 12.43 12.68 12.80 12.92 13.17 3586 1
## phi    0.14  0.00 0.06  0.04  0.10  0.14  0.18  0.26 3594 1
## sigma2  7.92  0.01 0.66  6.73  7.46  7.88  8.33  9.33 4058 1
## lp   -455.35  0.03 1.22 -458.63 -455.90 -455.02 -454.47 -453.96 2033 1
##
## Samples were drawn using NUTS(diag_e) at Tue May 9 10:30:51 2023.

```

```
## For each parameter, n_eff is a crude measure of effective sample size.  
## and Rhat is the potential scale reduction factor on split chains (at  
## convergence, Rhat=1).
```

```
print(fit_y)
```

```
## Inference for Stan model: d85eaf939f1fe064d5bdde49c5af5b23.  
## 4 chains, each with iter=2000; warmup=1000; thin=1;  
## post-warmup draws per chain=1000, total post-warmup draws=4000.  
##  
##      mean se_mean sd  2.5%  25%  50%  75%  97.5% n_eff Rhat  
## mu      10.04   1.21 17.77 -16.09  5.62  8.79 11.80  50.19  215 1.02  
## phi      0.96   0.00  0.02  0.91  0.94  0.96  0.97  1.00  584 1.01  
## sigma2    9.14   0.02  0.76  7.74  8.61  9.07  9.64 10.75 1608 1.00  
## lp     -476.53   0.06  1.48 -480.00 -477.38 -476.14 -475.37 -474.75  614 1.01  
##  
## Samples were drawn using NUTS(diag_e) at Tue May 9 10:30:56 2023.  
## For each parameter, n_eff is a crude measure of effective sample size.  
## and Rhat is the potential scale reduction factor on split chains (at  
## convergence, Rhat=1).
```

*#Conclusion: As the values in the credible interval are more narrow for the one where phi=0.2 were used, we get a better understanding of the true values for this one compared to the one where phi=0.95 were used. For this one, the values are really wide in the credible intervals*

```
##ii)  
##Joint posterior for X  
plot(post_x$mu[1000:2000],post_x$phi[1000:2000],xlab="My vals",ylab="Phi vals",  
main="Joint posterior of mu and phi (X)")
```

```
##Joint posterior for Y  
plot(post_y$mu[1000:2000],post_y$phi[1000:2000],xlab="My vals",ylab="Phi vals",  
main="Joint posterior of mu and phi (Y)")
```

*#Conclusion: We can distinguish a convergence for the first sampler. However it is more difficult to distinguish a convergence for the second sampler as shown in the figure. This follows the previous reasoning, that the second sampler have a wider span for its values compared to the first sampler*

## Appendix

### Exercise 1

```
data=readRDS("Precipitation.rds") #Reading data
data = log(data)

#a)

# Initial setting (prior settings)
my0 = 1
tao0=1
v0=1
sigma0sq=1
nDraws = 1000
n = nDraws
true_mean = mean(data)
vn = v0 + n

#Lists for storing draws
my_vals = rep(0,nDraws)
sigma_vals = rep(0,nDraws)
sigma_vals[1] = 1 #Initial sigmasq value to be used

#Drawing posteriors
for(i in 1:nDraws){

  w = (n/sigma_vals[i])/((n/sigma_vals[i])+(1/tao0**2)) #Calculating weights
  myn = w*true_mean+(1-w)*my0 #Calculating myn
  taonsq = 1/((n/sigma_vals[i])+(1/tao0)) #Calculating taonsq
  my_vals[i] = rnorm(1,mean=myn, taonsq) #Calculating posterior my_val

  if (i <= (nDraws-1)){
    chi_val = rchisq(1,vn) #Drawing random value from chi distribution
    sigmansq = (v0*sigma0sq+sum((data-my_vals[i])**2))/(n+v0) #Calculating approximation
    sigmasq
    sigma_vals[i+1] = vn*sigmansq/chi_val #Calculating posterior sigma_val
  }
}
```

#Plotting my convergence

```
plot(my_vals[10:nDraws], .main = "Plot of gibbs sampled distribution for my", type = "l")
abline(h=mean(my_vals), col="red") #Expected value
```

#Plotting sigma convergence

```
plot(sqrt(sigma_vals[10:nDraws]), .main = "Plot of gibbs sampled distribution for sigma", type = "l")
abline(h=mean(sqrt(sigma_vals)), col="red") #Expected value
```

#Calculating inefficiency factor (IF)

#Autocorrelation for my

a\_my = acf(my\_vals)

inefficiency\_factor\_my = 1+2\*sum(a\_my\$acf[-1])

inefficiency\_factor\_my

## [1] 0.5193648

#Autocorrelation for my

a\_sigma = acf(sqrt(sigma\_vals))

inefficiency\_factor\_sigma = 1+2\*sum(a\_sigma\$acf[-1])

inefficiency\_factor\_sigma

## [1] 1.04877

#Conclusion: As the plots suggest, we can't distinguish any convergence for the mean and standard deviation

#as we clearly can see that the values oscillate around the means (red line)

#Moreover, the inefficiency factors shows how correlated the draws are, as we can see, the results are

#relatively close to 1 which indicates low correlation

#b)

#1) Histogram or kernel density estimate of the daily preception y1.....yn

hist(data) #Histogram with true data

#2) The resulting posterior predictive density  $p(\tilde{y}|y_1, \dots, y_n)$  using the simulated posterior draws from (a)

predictions = rep(0,nDraws)

for (i in 1:nDraws){

predictions[i] = rnorm(1,my\_vals[i],sigma\_vals[i])

}

hist(predictions)

#Conclusion: As we can see, the draws from the predictions differ slightly from the true distribution.

#The main difference is that the normal distribution have a bit longer tails on each end, however

#the distributions seem to be centered around the same values in both plots

## Exercise 2

```
# Metropolis Random Walk for Poisson regression.
# Consider the following Poisson regression model  $y_i | \beta \sim \text{Poisson}[\exp(x_i \text{ transpose} \% \% \beta)]$ ,  $i=1, \dots, n$ 
# where  $y_i$  is the count for the  $i$ th observation in the sample
# and  $x_i$  is the  $p$ -dimensional vector with covariate observations for the  $i$ th observation.
# Use the data set eBayNumberOfBidderData.dat. This dataset contains observations from 1000 eBay auctions of coins.
# The response variable is nBids and records the number of bids in each auction.
# The remaining variables are features/covariates (x):
# Const (for the intercept)
# PowerSeller (equal to 1 if the seller is selling large volumes on eBay)
# VerifyID (equal to 1 if the seller is a verified seller by eBay)
# Sealed (equal to 1 if the coin was sold in an unopened envelope)
# MinBlem (equal to 1 if the coin has a minor defect)
# MajBlem (equal to 1 if the coin has a major defect)
# LargNeg (equal to 1 if the seller received a lot of negative feedback from customers)
# LogBook (logarithm of the book value of the auctioned coin according to expert sellers. Standardized)
# MinBidShare (ratio of the minimum selling price (starting price) to the book value. Standardized).

# (a) Obtain the maximum likelihood estimator of  $\beta$  in the Poisson regression model for the eBay data
# [Hint: glm.R, don't forget that glm() adds its own intercept so don't input the covariate Const].
# Which covariates are significant?

library(mvtnorm)

ebay_data <- (read.table("eBayNumberOfBidderData.dat", header=TRUE))
Xnames <- names(ebay_data[,2:ncol(ebay_data)]) #Names of columns
n_obs <- dim(ebay_data)[1] #number of observations
data <- ebay_data[, -2] #training and test data without Const
#y <- as.numeric(data[, 1]) #target output

# maximum likelihood estimator of beta
model = glm(formula=nBids~., data=data, family=poisson())
summary(model)

##
## Call:
## glm(formula = nBids ~ ., family = poisson(), data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.5800 -0.7222 -0.0441  0.5269  2.4605
```

```
##
## Coefficients:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept)  1.07244      0.03077  34.848 < 2e-16 ***
## PowerSeller -0.02054      0.03678  -0.558  0.5765
## VerifyID    -0.39452      0.09243  -4.268 1.97e-05 ***
## Sealed      0.44384      0.05056   8.778 < 2e-16 ***
## Minblem    -0.05220      0.06020  -0.867  0.3859
## MajBlem    -0.22087      0.09144  -2.416  0.0157 *
## LargNeg    0.07067      0.05633   1.255  0.2096
## LogBook    -0.12068      0.02896  -4.166 3.09e-05 ***
## MinBidShare -1.89410      0.07124 -26.588 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 2151.28  on 999  degrees of freedom
## Residual deviance:  867.47  on 991  degrees of freedom
## AIC: 3610.3
##
## Number of Fisher Scoring iterations: 5
```

#the z-value shows that VerifyID, Sealed, MajBlem, LogBook and MinBidShare are significant

# (b) Let's do a Bayesian analysis of the Poisson regression. Let the prior be  $\beta \sim N[0, 100 \cdot (X^T X)^{-1}]$   
# where  $X$  is the  $n \times p$  covariate matrix. This is a commonly used prior, which is called Zellner's g-prior.  
# Assume first that the posterior density is approximately multivariate normal:  
#  $\beta | y \sim N(\beta_{\text{tilde}}, J_y^{-1}(\beta_{\text{tilde}}))$   
# where  $\beta_{\text{tilde}}$  is the posterior mode and  $J_y(\beta_{\text{tilde}})$  is the negative Hessian at the posterior mode.  
#  $\beta_{\text{tilde}}$  and  $J_y(\beta_{\text{tilde}})$  can be obtained by numerical optimization (optim.R)  
# exactly like you already did for the logistic regression in Lab 2  
# (but with the log posterior function replaced by the corresponding one for the Poisson model,  
# which you have to code up.).

```
#function for Poisson logPosterior
logPostPoisson = function(beta,y,X,mu,Sigma) {
  beta <- matrix(beta, nrow=1)
  linPred <- X%*%t(beta)
  logLik <- sum( - exp(linPred) + linPred*y - log(factorial(y)) ) #loglikelihood
  logPrior <- dmvnorm(beta, mu, Sigma, log=TRUE); #prior
  return(logLik + logPrior) #prior and likelihood are added since we use log
}
```



```

#initialize variables
y <- as.numeric(ebay_data[,1]) #target output
X <- as.matrix(ebay_data[, -1]) #training data
X <- apply(X, c(1, 2), as.numeric) #convert from characters to doubles
n_col <- dim(X)[2] #number of columns in training data
initVal <- rep(0, n_col) #initial values for Beta
mu <- as.matrix(rep(0, n_col)) # Prior mean vector for Beta
Sigma <- 100*solve(t(X)%*%X) # Prior covariance matrix for Beta
logPost <- logPostPoisson; #function to be minimized

set.seed(123)
OptimRes <-
optim(initVal, logPost, gr=NULL, y=X, mu, Sigma, method=c("BFGS"), control=list(fnscale=-1), hessian=TRUE)

# Printing the results of mean and variance of normally distributed Beta
OptimBeta <- OptimRes$par
names(OptimBeta) <- Xnames # Naming the coefficient by covariates
PostCov <- solve(-OptimRes$hessian) #get covariance matrix by inverting the negative
hessian
approxPostStd <- sqrt(diag(PostCov)) # Computing approximate standard deviations.
names(approxPostStd) <- Xnames # Naming the coefficient by covariates
print('Posterior beta values:')

## [1] "Posterior beta values:"

print(OptimBeta)

##      Const PowerSeller   VerifyID      Sealed      Minblem      MajBlem
## 1.06984118 -0.02051246 -0.39300599 0.44355549 -0.05246627 -0.22123840
##      LargNeg      LogBook MinBidShare
## 0.07069683 -0.12021767 -1.89198501

print('Posterior standard deviation:')

## [1] "Posterior standard deviation:"

print(approxPostStd)

##      Const PowerSeller   VerifyID      Sealed      Minblem      MajBlem
## 0.03074837 0.03678418 0.09227871 0.05057448 0.06020470 0.09146070
##      LargNeg      LogBook MinBidShare
## 0.05634767 0.02895635 0.07109682

#the result seems reasonable since they are similar to a)

# (c) Let's simulate from the actual posterior of beta using the Metropolis algorithm
# and compare the results with the approximate results in b).

```

# Program a general function that uses the Metropolis algorithm to generate random draws from an  
# arbitrary posterior density. In order to show that it is a general function for  
# any model, we denote the vector of model parameters by  $\theta$ . Let the proposal density  
# be the multivariate normal density mentioned in Lecture 8 (random walk Metropolis):  
#  $\theta_p|\theta(i-1) \sim N(\theta(i-1), c \cdot \Sigma)$  where  $\Sigma = J\gamma^{-1}(\beta^*)$  was obtained in b).  
# The value  $c$  is a tuning parameter and should be an input to your Metropolis function.  
# The user of your Metropolis function should be able to supply her own posterior density  
function, not  
# necessarily for the Poisson regression, and still be able to use your Metropolis  
# function. This is not so straightforward, unless you have come across function  
# objects in R. The note HowToCodeRWM.pdf in Lisam describes how you can do this in R.

#Now, use your new Metropolis function to sample from the posterior of  $\beta$   
# in the Poisson regression for the eBay dataset. Assess MCMC convergence by graphical  
methods.

```
RWMSampler <- function(oldTheta, logPostFunc, c, PostCov, ... ) {
  proposedTheta <- rmvnorm(1, mean=oldTheta, sigma=c*PostCov)
  alpha <- min(1, exp(logPostFunc(proposedTheta, ...) - logPostFunc(oldTheta, ...)))
  t <- runif(1)
  if (alpha>t) {
    return(list(proposedTheta, alpha))
  } else {
    return(list(oldTheta, alpha))
  }
}
```

```
nDraws <- 4000
```

```
c <- 0.7
```

```
beta <- matrix(0, nDraws, n_col) #matrix to fill with samples from RWM
```

```
alpha_vector <- rep(0, nDraws) #will be used to calculate mean alpha
```

```
logPost <- logPostPoisson
```

```
for(i in 1:nDraws) {
  output <- RWMSampler(beta[i,], logPost, c, PostCov, y, X, mu, Sigma)
  alpha_vector[i] <- output[[2]]
  if(i<nDraws) {
    beta[i+1,] <- output[[1]]
  }
}
```

```
iterations=seq(1,nDraws,1)
```

```
par(mfrow=c(3,3))
```

```
for (i in 1:n_col) {
```

```
  plot(iterations, beta[,i], type="l", main=paste("MCMC Convergence", Xnames[i]),
```

```

      ylab=Xnames[i])
    }

par(mfrow=c(1,1), new=FALSE)

# Calculate average alpha
average_alpha <- mean(alpha_vector)
average_alpha #0.2531, has average acceptance probability between 25-30%

## [1] 0.2621183

# (d) Use the MCMC draws from c) to simulate from the predictive distribution of
# the number of bidders in a new auction with the characteristics below.
#Plot the predictive distribution. What is the probability of no bidders in this new auction?
# PowerSeller = 1
# VerifyID = 0
# Sealed = 1
# MinBlem = 0
# MajBlem = 1
# LargNeg = 0
# LogBook = 1.2
# MinBidShare = 0.8

nDraws <- 10000
x <- c(1,1,0,1,0,1,0,1,2,0,8)
post_beta <- beta[(1000:nrow(beta)),] #select samples after convergence
mean <- exp(post_beta%*%x) #vector of means
samples <- rpois(nDraws, mean)
barplot(table(samples), main="Histogram of number of bidders")

prob_no_bidders <- sum(samples==0)/sum(samples) #probability of no bidders
print(prob_no_bidders)

## [1] 0.739524

#the probability of no bidders is 74%

```

### Exercise 3

#a)

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```
## Loading required package: ggplot2
```

```
## rstan (Version 2.21.8, GitRev: 2e1f913d3ca3)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
```

```
## options(mc.cores = parallel::detectCores()).
```

```
## To avoid recompilation of unchanged Stan programs, we recommend calling
```

```
## rstan_options(auto_write = TRUE)
```

# Initial values

```
my=13
```

```
sigmasq=3
```

```
T=300
```

```
x1 = my
```

```
phi = seq(-1,1,0.25)
```

#Autoregression function

```
ar_function = function(my,sigmasq,T,x1,phi){
```

```
  result = rep(0,T)
```

```
  result[1] = x1
```

```
  for(i in 2:T){
```

```
    epsilon = rnorm(1,0,sigmasq)
```

```
    result[i] = my + phi*(result[i-1]-my) + epsilon
```

```
  }
```

```
  return(result)
```

```
}
```

#Calculating output

```
result = matrix(0,T,length(phi))
```

```
for(i in 1:length(phi)){
```

```
  result[,i] = ar_function(my,sigmasq,T,x1,phi[i])
```

```
}
```

#Plotting result for phi=-1 and phi=1

```
plot(seq(1,300,1),result[,1],xlab="Order",ylab="Realization",main="Realization for phi =  
-1",type="l")
```

```
plot(seq(1,300,1),result[,3],xlab="Order",ylab="Realization",main="Realization for phi =  
-0.5",type="l")
```

```
plot(seq(1,300,1),result[,7],xlab="Order",ylab="Realization",main="Realization for phi =  
0.5",type="l")
```

```
plot(seq(1,300,1),result[,9],xlab="Order",ylab="Realization",main="Realization for phi = 1",type="l")
```

*#Conclusion: Increased phi values reduces oscillations. We can see clear oscillations for phi=-1*

*#and that the output becomes more and more correlated with increased phi values*

*#b)*

*#Drawing x1:T values using phi=0.2*

```
x=rep(0,T)
```

```
x=ar_function(my,sigma2,T,x1,0.2)
```

*#Drawing y1:T values using phi=0.95*

```
y=rep(0,T)
```

```
y=ar_function(my,sigma2,T,x1,0.95)
```

*#Implementing Stan model*

```
StanModel = '
```

```
data {
```

```
  int<lower=0> N;
```

```
  vector[N] z;
```

```
}
```

```
parameters {
```

```
  real mu;
```

```
  real phi;
```

```
  real<lower=0> sigma2;
```

```
}
```

```
model {
```

```
  for(i in 2:N){
```

```
    z[i] ~ normal(mu+phi*(z[i-1]-mu),sqrt(sigma2));
```

```
  }
```

```
}
```

*#For X*

```
data <- list(N=T, z=x)
```

```
fit_x <- stan(model_code=StanModel,data=data)
```

*#For Y*

```
data <- list(N=T, z=y)
```

```
fit_y <- stan(model_code=StanModel,data=data)
```

*#Extract posterior draws*

```
post_x = extract(fit_x)
```

```
post_y = extract(fit_y)
```

*#i)*

*#AR-process for X*

*#Calculating mean and credible interval for all parameters*

*#Can be achieved with print(fit\_x) too*

```
mean_mu = mean(post_x$mu)
cred_int_mu = quantile(post_x$mu, probs=c(0.025,0.975))
print(mean_mu)
```

```
## [1] 12.79963
```

```
print(cred_int_mu)
```

```
## 2.5% 97.5%
```

```
## 12.43112 13.17352
```

```
mean_phi = mean(post_x$phi)
cred_int_phi = quantile(post_x$phi, probs=c(0.025,0.975))
print(mean_phi)
```

```
## [1] 0.1432679
```

```
print(cred_int_phi)
```

```
## 2.5% 97.5%
```

```
## 0.03507495 0.25712395
```

```
mean_sigma2 = mean(post_x$sigma2)
cred_int_sigma2 = quantile(post_x$sigma2, probs=c(0.025,0.975))
print(mean_sigma2)
```

```
## [1] 7.918592
```

```
print(cred_int_sigma2)
```

```
## 2.5% 97.5%
```

```
## 6.726024 9.330297
```

*#AR-process for Y*

*#Calculating mean and credible interval for all parameters*

```
mean_mu = mean(post_y$mu)
cred_int_mu = quantile(post_y$mu, probs=c(0.025,0.975))
print(mean_mu)
```

```
## [1] 10.04247
```

```
print(cred_int_mu)
```

```
## 2.5% 97.5%
```

```
## -16.08935 50.19349
```

```
mean_phi = mean(post_y$phi)
cred_int_phi = quantile(post_y$phi, probs=c(0.025,0.975))
print(mean_phi)
```

```
## [1] 0.957765
```

```

print(cred_int_phi)

##      2.5%      97.5%
## 0.9116279 1.0007403

mean_sigma2 = mean(post_y$sigma2)
cred_int_sigma2 = quantile(post_y$sigma2, probs=c(0.025,0.975))
print(mean_sigma2)

## [1] 9.138192

print(cred_int_sigma2)

##      2.5%      97.5%
## 7.744661 10.745055

#Plotting using print
print(fit_x)

## Inference for Stan model: d85eaf939f1fe064d5bdde49c5af5b23.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##      mean se mean sd  2.5%  25%  50%  75%  97.5% n_eff Rhat
## mu      12.80  0.00 0.19 12.43 12.68 12.80 12.92 13.17 3586 1
## phi      0.14  0.00 0.06  0.04  0.10  0.14  0.18  0.26 3594 1
## sigma2    7.92  0.01 0.66  6.73  7.46  7.88  8.33  9.33 4058 1
## lp     -455.35  0.03 1.22 -458.63 -455.90 -455.02 -454.47 -453.96 2033 1
##
## Samples were drawn using NUTS(diag_e) at Tue May 9 10:30:51 2023.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

print(fit_y)

## Inference for Stan model: d85eaf939f1fe064d5bdde49c5af5b23.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##      mean se mean sd  2.5%  25%  50%  75%  97.5% n_eff Rhat
## mu      10.04  1.21 17.77 -16.09  5.62  8.79 11.80 50.19 215 1.02
## phi      0.96  0.00 0.02  0.91  0.94  0.96  0.97  1.00 584 1.01
## sigma2    9.14  0.02 0.76  7.74  8.61  9.07  9.64 10.75 1608 1.00
## lp     -476.53  0.06 1.48 -480.00 -477.38 -476.14 -475.37 -474.75 614 1.01
##
## Samples were drawn using NUTS(diag_e) at Tue May 9 10:30:56 2023.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```

#Conclusion: As the values in the credible interval are more narrow for the one where  $\phi=0.2$   
#were used, we get a better understanding of the true values for this one compared to the one where  
# $\phi=0.95$  were used. For this one, the values are really wide in the credible intervals

#ii)

#Joint posterior for X

plot(post\_x\$mu[1000:2000],post\_x\$phi[1000:2000],xlab="My vals",ylab="Phi vals",  
main="Joint posterior of mu and phi (X)")

#Joint posterior for Y

plot(post\_y\$mu[1000:2000],post\_y\$phi[1000:2000],xlab="My vals",ylab="Phi vals",  
main="Joint posterior of mu and phi (Y)")

#Conclusion: We can distinguish a convergence for the first sampler. However it is more difficult to distinguish a

#convergence for the second sampler as shown in the figure. This follows the previous reasoning, that the

#second sampler have a wider span for its values compared to the first sampler